

# Work Sample

J. Antonio García

09/11/2020

## Reduce maintenance cost through predictive techniques

### Preamble

```
# html_notebook

packs <- c('lubridate', 'dplyr', 'ggplot2', 'caret', 'MLmetrics', 'gbm', 'e1071', 'LiblineaR',
          'xgboost', 'randomForest', 'doParallel')
index <- packs %in% row.names(installed.packages())
if (any(!index)){
  sapply(packs[!index], FUN=install.packages )
}

require(lubridate) # easy handle datetimes
require(dplyr) # like SQL in R, and also load pipe operator
require(ggplot2) # easy, fast and nice plots
require(caret) # a toolbox
require(MLmetrics) # metric in one line
require(gbm) # boosting alg.
require(e1071) # numerical routines for svm implementation
require(LiblineaR) # another implementation for RLogReg
require(xgboost) # The implementation
require(randomForest) # nice RF implementation
require(doParallel) #ugly parallel in R but useful

cl <- makePSOCKcluster(2)
registerDoParallel(cl)
```

Set up parallel enviroment

### EDA

```
rm(list=ls()) # clean env.
t1 <- Sys.time()
data.raw <- read.csv(file='device_failure.csv') # few records kernels function works fine
print(sum(is.na(data.raw))) # NOT NULLS! THANKS A LOT :D

## [1] 0

data.raw %>% arrange(device, date) %>% mutate(date = ymd(date) ) -> data.raw
data.raw %>% group_by(device) %>% arrange(device, date) %>%
```

```
mutate( l.attribute1 = lag(attribute1),
        l.attribute2 = lag(attribute2),
        l.attribute3 = lag(attribute3),
        l.attribute4 = lag(attribute4),
        l.attribute5 = lag(attribute5),
        l.attribute6 = lag(attribute6),
        l.attribute7 = lag(attribute7),
        l.attribute8 = lag(attribute8),
        l.attribute9 = lag(attribute9)) -> data.raw
data.raw <- na.omit(data.raw)
```

Each device has 0 or 1 failure, and if has a failure it's the last row

```
n.fails.index <- which(data.raw$failure==1) #only 106 failures
nn.fails <- data.raw[ rep(n.fails.index, each=9) + -4:4, ]
head(nn.fails, 50 )
```

```
## # A tibble: 50 x 21
## # Groups:   device [11]
##   date      device failure attribute1 attribute2 attribute3 attribute4
##   <date>    <chr>    <int>      <int>      <int>      <int>      <int>
## 1 2015-01-15 S1F02~      0  222474632      0      0      1
## 2 2015-01-16 S1F02~      0  243825496      0      0      1
## 3 2015-01-17 S1F02~      0  20761856      0      0      1
## 4 2015-01-18 S1F02~      0  41291000      0      0      1
## 5 2015-01-19 S1F02~      1  64499464      0      0      1
## 6 2015-01-02 S1F02~      0  63705712      0      1      0
## 7 2015-01-03 S1F02~      0  53868456      0      1      0
## 8 2015-01-04 S1F02~      0  4263992      0      1      0
## 9 2015-01-05 S1F02~      0  37773128      0      1      0
## 10 2015-07-30 S1F03~      0  3869656      232      0      0
## # ... with 40 more rows, and 14 more variables: attribute5 <int>,
## #   attribute6 <int>, attribute7 <int>, attribute8 <int>, attribute9 <int>,
## #   l.attribute1 <int>, l.attribute2 <int>, l.attribute3 <int>,
## #   l.attribute4 <int>, l.attribute5 <int>, l.attribute6 <int>,
## #   l.attribute7 <int>, l.attribute8 <int>, l.attribute9 <int>
```

And in general, the devices present at most one fault and from which no information is recorded about them. Above all we are in a case where the variable to predict has a **strong positive bias**, more than 99% of the records are not failures, a very common case in practice ...

```
data.raw %>% filter(failure==1) %>% group_by(device) %>% summarise(n=n()) -> t
summary(t$n)
```

```
##   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##      1      1      1      1      1      1
```

```
table(data.raw$failure) / dim(data.raw)[1]
```

```
##
##           0           1
## 0.9991404825 0.0008595175
```

```
summary(data.raw)
```


```
##      date      device      failure
##   Min.   :2015-01-02   Length:123325   Min.    :0.0000000
```

```
## 1st Qu.:2015-02-10   Class :character   1st Qu.:0.0000000
## Median :2015-03-28   Mode  :character   Median :0.0000000
## Mean   :2015-04-17                   Mean   :0.0008595
## 3rd Qu.:2015-06-18                   3rd Qu.:0.0000000
## Max.   :2015-11-02                   Max.    :1.0000000
## attribute1          attribute2          attribute3          attribute4
## Min.   :      0   Min.   :      0.0   Min.   :      0.00   Min.   :      0.000
## 1st Qu.: 61311432   1st Qu.:      0.0   1st Qu.:      0.00   1st Qu.:      0.000
## Median :122785544   Median :      0.0   Median :      0.00   Median :      0.000
## Mean   :122391362   Mean   : 157.7   Mean   :      9.76   Mean   :      1.723
## 3rd Qu.:183330360   3rd Qu.:      0.0   3rd Qu.:      0.00   3rd Qu.:      0.000
## Max.   :244140480   Max.   :64968.0   Max.   :24929.00   Max.   :1666.000
## attribute5          attribute6          attribute7          attribute8
## Min.   : 1.00   Min.   :      8   Min.   : 0.0000   Min.   : 0.0000
## 1st Qu.: 8.00   1st Qu.:221528   1st Qu.: 0.0000   1st Qu.: 0.0000
## Median :10.00   Median :250060   Median : 0.0000   Median : 0.0000
## Mean   :14.24   Mean   :260377   Mean   : 0.2892   Mean   : 0.2892
## 3rd Qu.:12.00   3rd Qu.:310396   3rd Qu.: 0.0000   3rd Qu.: 0.0000
## Max.   :98.00   Max.   :689161   Max.   :832.0000   Max.   :832.0000
## attribute9          l.attribute1          l.attribute2          l.attribute3
## Min.   : 0.00   Min.   :      0   Min.   : 0.0   Min.   : 0.000
## 1st Qu.: 0.00   1st Qu.: 61298592   1st Qu.: 0.0   1st Qu.: 0.000
## Median : 0.00   Median :122798936   Median : 0.0   Median : 0.000
## Mean   : 12.11   Mean   :122390254   Mean   : 152.7   Mean   : 9.739
## 3rd Qu.: 0.00   3rd Qu.:183314520   3rd Qu.: 0.0   3rd Qu.: 0.000
## Max.   :18701.00   Max.   :244140480   Max.   :64968.0   Max.   :24929.000
## l.attribute4          l.attribute5          l.attribute6          l.attribute7
## Min.   : 0.000   Min.   : 1.00   Min.   :      8   Min.   : 0.0000
## 1st Qu.: 0.000   1st Qu.: 8.00   1st Qu.:221462   1st Qu.: 0.0000
## Median : 0.000   Median :10.00   Median :249721   Median : 0.0000
## Mean   : 1.665   Mean   :14.24   Mean   :260080   Mean   : 0.2532
## 3rd Qu.: 0.000   3rd Qu.:12.00   3rd Qu.:310207   3rd Qu.: 0.0000
## Max.   :1666.000   Max.   :98.00   Max.   :689062   Max.   :832.0000
## l.attribute8          l.attribute9
## Min.   : 0.0000   Min.   : 0.0
## 1st Qu.: 0.0000   1st Qu.: 0.0
## Median : 0.0000   Median : 0.0
## Mean   : 0.2532   Mean   : 12.1
## 3rd Qu.: 0.0000   3rd Qu.: 0.0
## Max.   :832.0000   Max.   :18701.0
```


## Focus on success stories

So we decided to focus on success stories to infer from them insights that allow us to carry out the task.


```
device.with.failures <- unique(data.raw$device[n.failures.index] )
data.sample <- data.raw[ data.raw$device %in% device.with.failures, ]
ggplot(data.sample, aes(attribute1, fill = as.character(failure), alpha=.01)) +
  geom_density() + theme_minimal() + theme(legend.position="none")
```



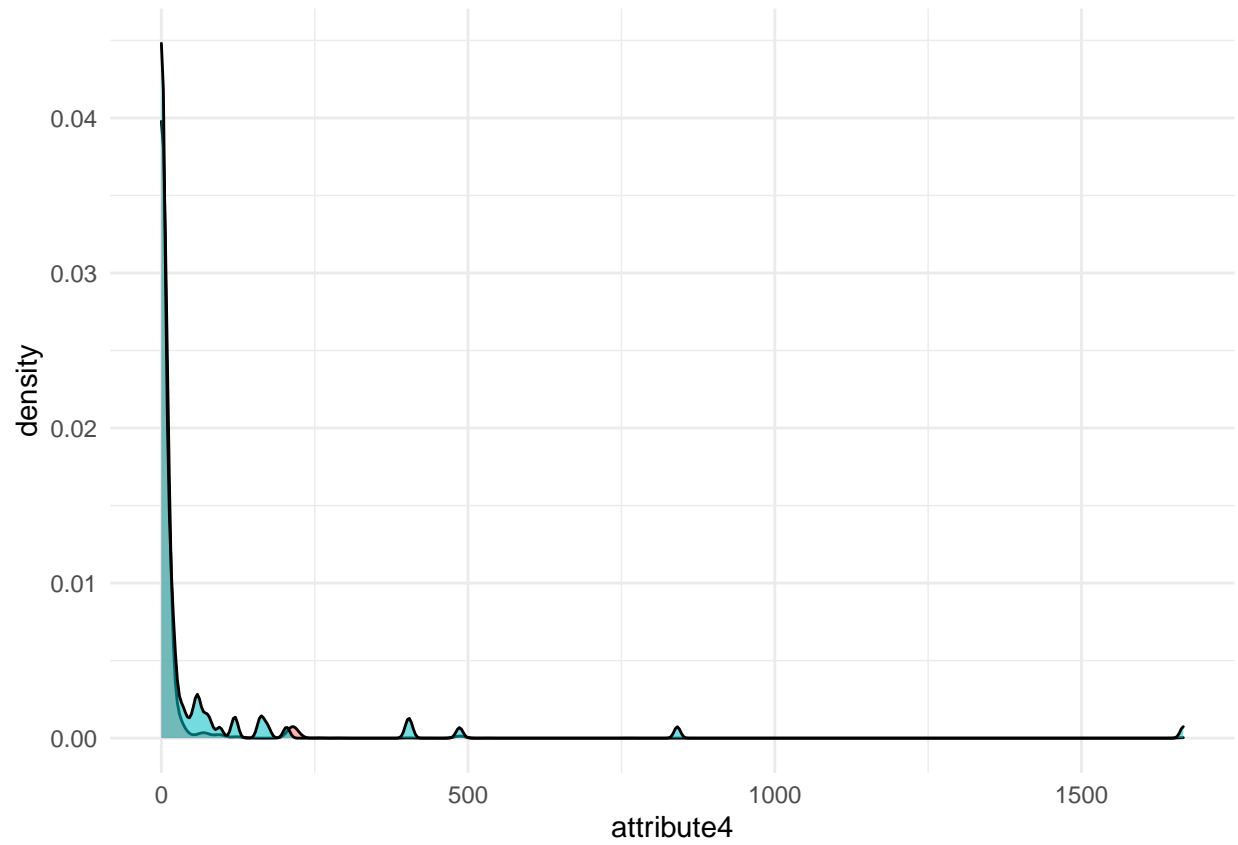
```
ggplot(data.sample, aes(attribute2, fill = as.character(failure), alpha=.01)) +  
  geom_density() + theme_minimal() + theme(legend.position="none")
```




```
ggplot(data.sample, aes(attribute3, fill = as.character(failure), alpha=.01)) +  
  geom_density() + theme_minimal() + theme(legend.position="none")
```



```
ggplot(data.sample, aes(attribute4, fill = as.character(failure), alpha=.01)) +  
  geom_density() + theme_minimal() + theme(legend.position="none")
```

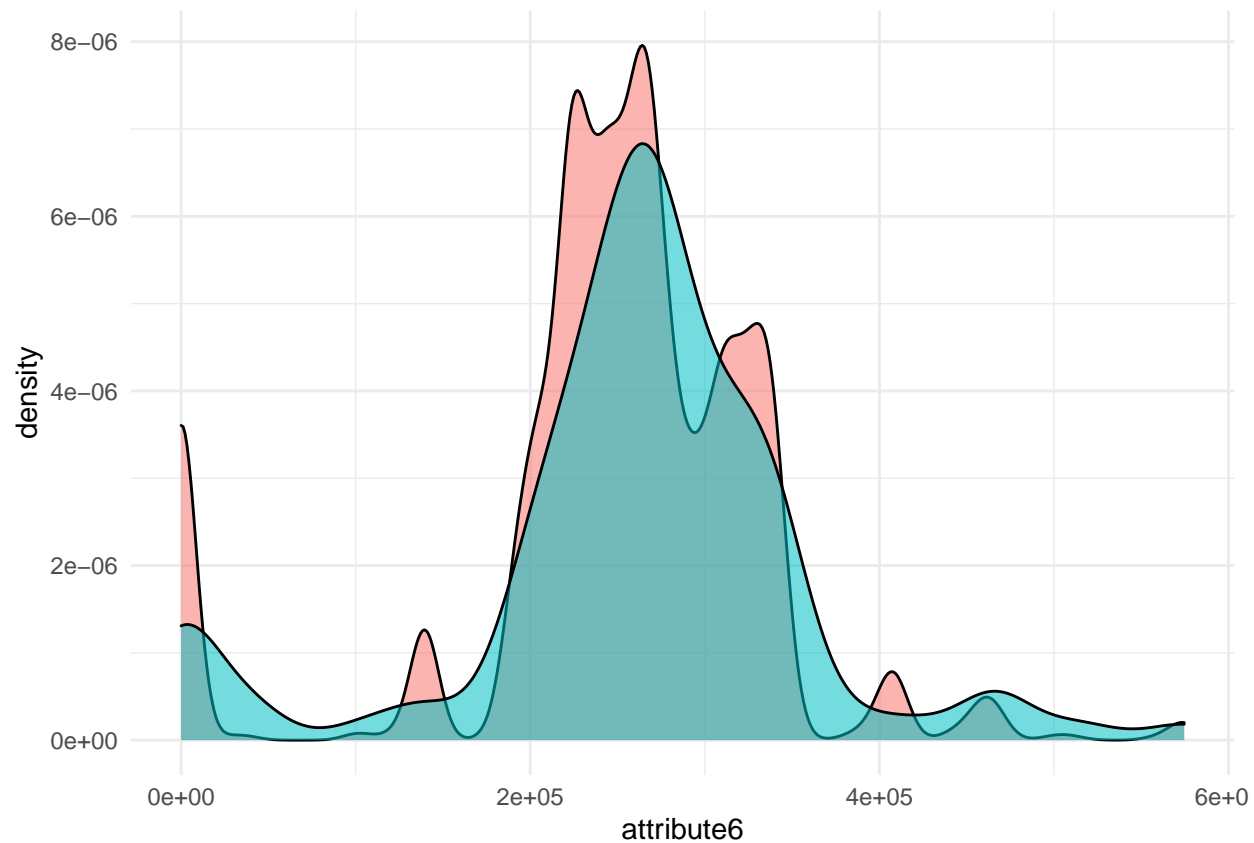


```
ggplot(data.sample, aes(attribute5, fill = as.character(failure), alpha=.01)) +  
  geom_density() + theme_minimal()+theme(legend.position="none")
```




```
ggplot(data.sample, aes(attribute6, fill = as.character(failure), alpha=.01)) +  
  geom_density() + theme_minimal() + theme(legend.position="none")
```

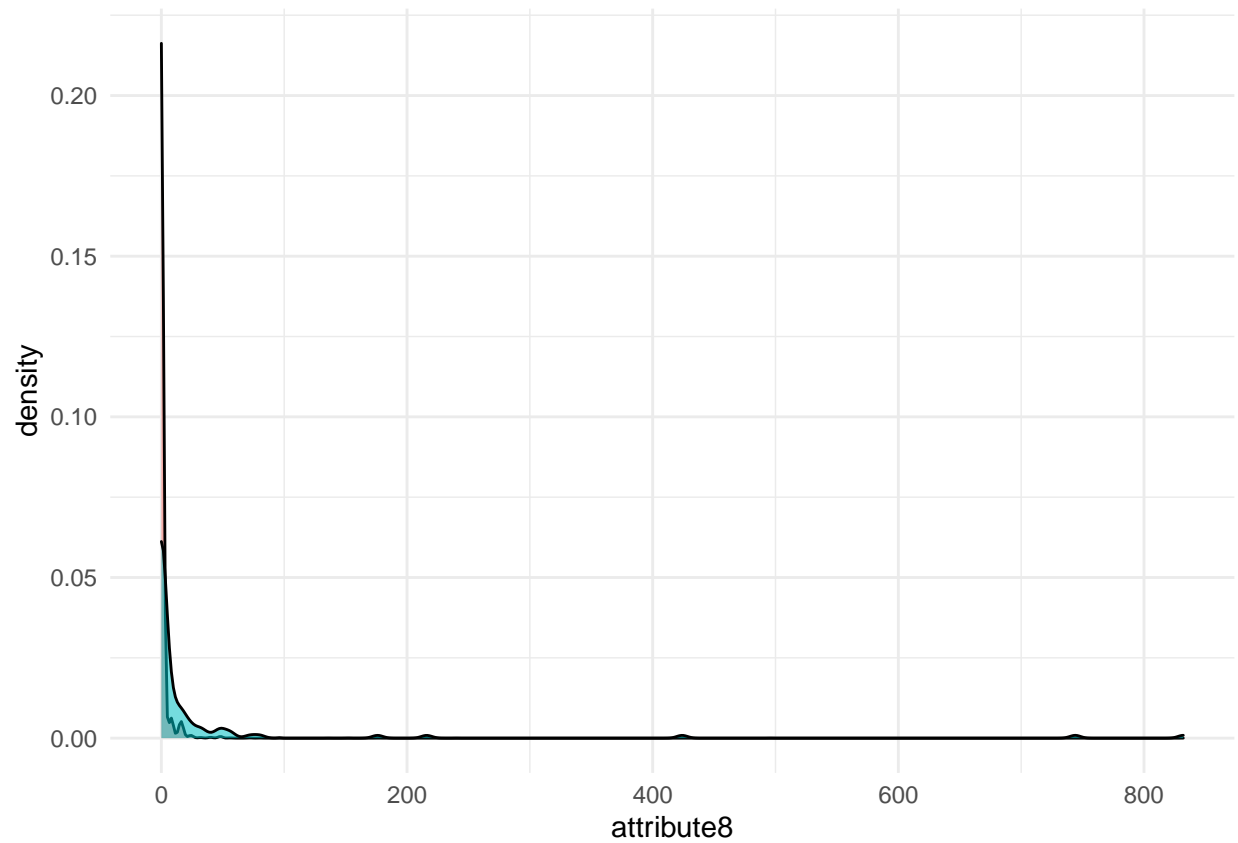





```
ggplot(data.sample, aes(attribute7, fill = as.character(failure), alpha=.01)) +  
  geom_density() + theme_minimal()+theme(legend.position="none")
```



```
ggplot(data.sample, aes(attribute8, fill = as.character(failure), alpha=.01)) +  
  geom_density() + theme_minimal() + theme(legend.position="none")
```



```
ggplot(data.sample, aes(attribute9, fill = as.character(failure), alpha=.01)) +  
  geom_density() + theme_minimal() + theme(legend.position="none")
```




Due to the distribution of some of the variables, we apply a non-linear transformation that allows us to more easily discriminate between failures and non-failures.

```
index.columns <- c(2, 3, 4, 7, 8, 9:18) + 3
# log features selected
data.sample[, names(data.sample)[index.columns]] <-
  log(data.sample[, names(data.sample)[index.columns]] + 1 )
# standar features
index.columns <- grep('attr', names(data.sample))
summary(data.sample)
```


```
##      date          device      failure      attribute1
## Min.   :2015-01-02   Length:10607   Min.    :0.000000   Min.    : 4224
## 1st Qu.:2015-01-31   Class :character   1st Qu.:0.000000   1st Qu.: 61045068
## Median :2015-03-12   Mode  :character   Median :0.000000   Median :123450384
## Mean   :2015-03-22                      Mean  :0.009993   Mean   :122672510
## 3rd Qu.:2015-05-02                      3rd Qu.:0.000000   3rd Qu.:184160428
## Max.   :2015-10-26                      Max.   :1.000000   Max.   :244135688
## attribute2      attribute3      attribute4      attribute5
## Min.   : 0.000   Min.   :0.0000   Min.   :0.0000   Min.   : 2.0
## 1st Qu.: 0.000   1st Qu.:0.0000   1st Qu.:0.0000   1st Qu.: 8.0
## Median : 0.000   Median :0.0000   Median :0.0000   Median : 9.0
## Mean   : 1.373   Mean   :0.2365   Mean   :0.5074   Mean   :14.2
## 3rd Qu.: 0.000   3rd Qu.:0.0000   3rd Qu.:0.0000   3rd Qu.:13.0
## Max.   :11.082   Max.   :5.7652   Max.   :7.4188   Max.   :91.0
## attribute6      attribute7      attribute8      attribute9
## Min.   : 19     Min.   :0.0000   Min.   :0.0000   Min.   :0.0000
```

```
## 1st Qu.:222560 1st Qu.:0.0000 1st Qu.:0.0000 1st Qu.:0.0000
## Median :256236 Median :0.0000 Median :0.0000 Median :0.0000
## Mean :248158 Mean :0.1895 Mean :0.1895 Mean :0.4359
## 3rd Qu.:299202 3rd Qu.:0.0000 3rd Qu.:0.0000 3rd Qu.:0.0000
## Max. :574599 Max. :6.7250 Max. :6.7250 Max. :7.0613
## l.attribute1 l.attribute2 l.attribute3 l.attribute4
## Min. : 8.349 Min. : 0.000 Min. :0.0000 Min. :0.0000
## 1st Qu.:17.925 1st Qu.: 0.000 1st Qu.:0.0000 1st Qu.:0.0000
## Median :18.629 Median : 0.000 Median :0.0000 Median :0.0000
## Mean :18.314 Mean : 1.353 Mean :0.2365 Mean :0.4941
## 3rd Qu.:19.031 3rd Qu.: 0.000 3rd Qu.:0.0000 3rd Qu.:0.0000
## Max. :19.313 Max. :11.082 Max. :5.7652 Max. :7.4188
## l.attribute5 l.attribute6 l.attribute7 l.attribute8
## Min. :1.099 Min. : 2.996 Min. :0.0000 Min. :0.0000
## 1st Qu.:2.197 1st Qu.:12.312 1st Qu.:0.0000 1st Qu.:0.0000
## Median :2.303 Median :12.453 Median :0.0000 Median :0.0000
## Mean :2.458 Mean :11.833 Mean :0.1794 Mean :0.1794
## 3rd Qu.:2.639 3rd Qu.:12.609 3rd Qu.:0.0000 3rd Qu.:0.0000
## Max. :4.522 Max. :13.261 Max. :6.7250 Max. :6.7250
## l.attribute9
## Min. :0.000
## 1st Qu.:0.000
## Median :0.000
## Mean :0.435
## 3rd Qu.:0.000
## Max. :7.061
```


```
for ( i in index.columns ){
  temp <- data.sample[, names(data.sample)[i]]
  data.sample[, names(data.sample)[i]] <- scale(temp)
}
ggplot(data.sample, aes(attribute1, fill = as.character(failure), alpha=.01)) +
  geom_density() + theme_minimal() + theme(legend.position="none")
```



```
ggplot(data.sample, aes(attribute2, fill = as.character(failure), alpha=.01)) +  
  geom_density() + theme_minimal() + theme(legend.position="none")
```

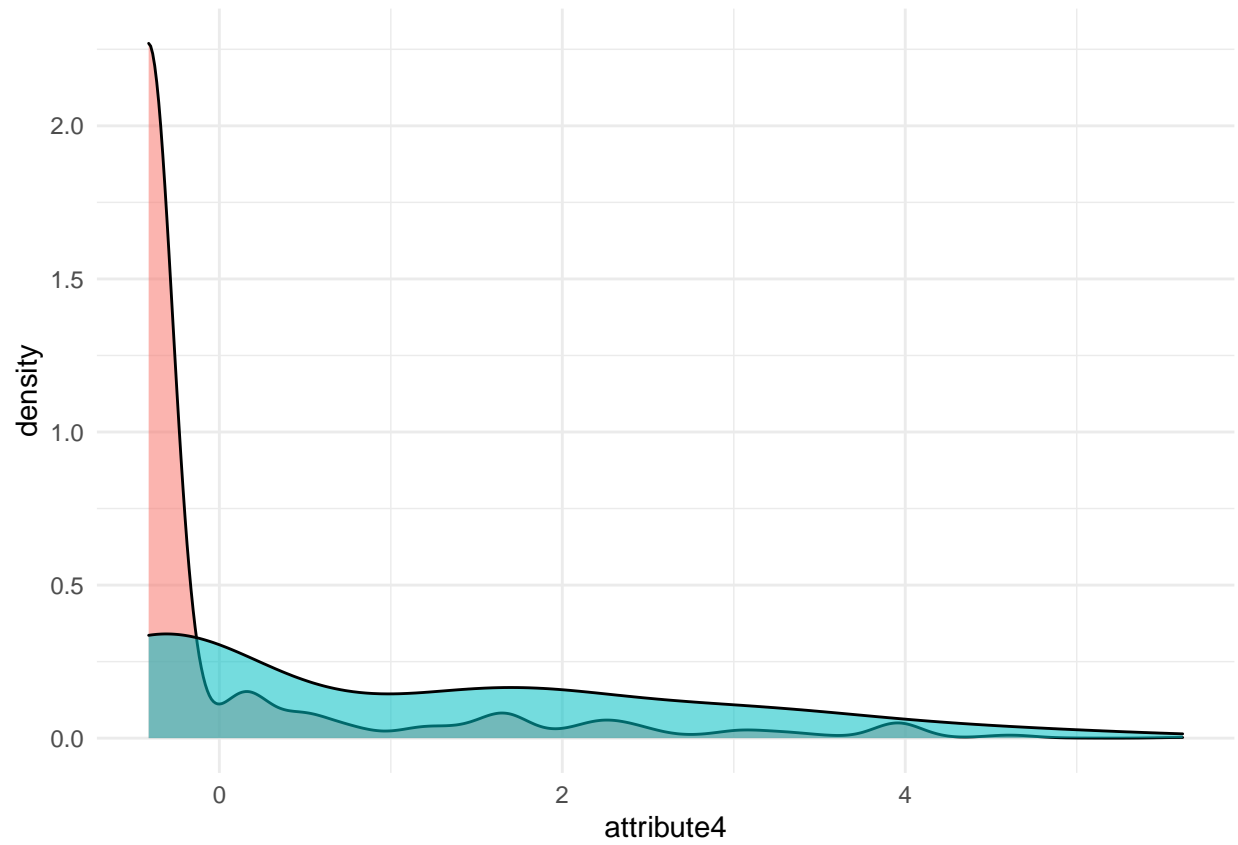


```
ggplot(data.sample, aes(attribute3, fill = as.character(failure), alpha=.01)) +  
  geom_density() + theme_minimal() + theme(legend.position="none")
```




```
ggplot(data.sample, aes(attribute4, fill = as.character(failure), alpha=.01)) +  
  geom_density() + theme_minimal() + theme(legend.position="none")
```







```
ggplot(data.sample, aes(attribute5, fill = as.character(failure), alpha=.01)) +  
  geom_density() + theme_minimal() + theme(legend.position="none")
```




```
ggplot(data.sample, aes(attribute6, fill = as.character(failure), alpha=.01)) +  
  geom_density() + theme_minimal() + theme(legend.position="none")
```



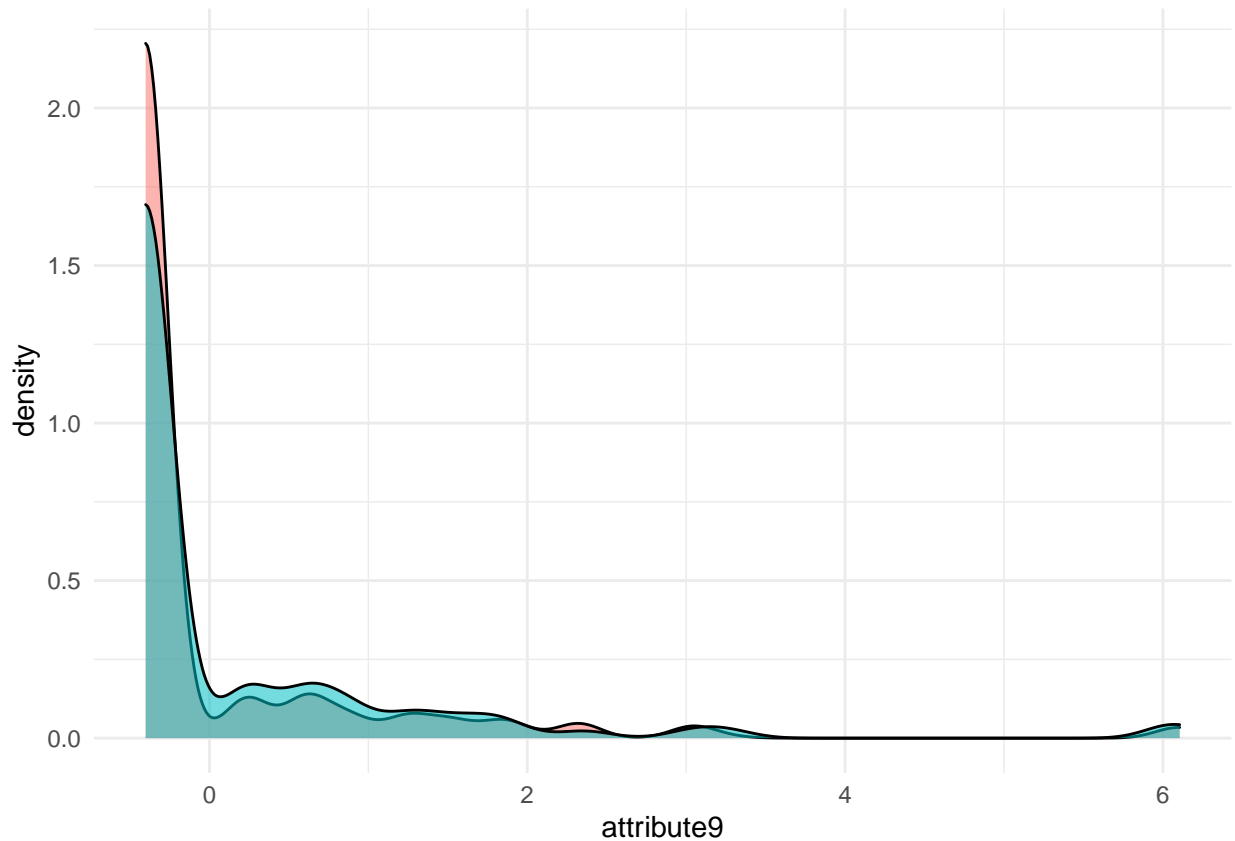
```
ggplot(data.sample, aes(attribute7, fill = as.character(failure), alpha=.01)) +  
  geom_density() + theme_minimal() + theme(legend.position="none")
```



```
ggplot(data.sample, aes(attribute8, fill = as.character(failure), alpha=.01)) +  
  geom_density() + theme_minimal() + theme(legend.position="none")
```



```
ggplot(data.sample, aes(attribute9, fill = as.character(failure), alpha=.01)) +  
  geom_density() + theme_minimal() + theme(legend.position="none")
```



After we divided our sample to continue with a preselection of models, among the enormous variety of algorithms and implementations that exist, we decided to report 4, because they are interpretable models and easy to explain to non-specialized people.

Since we are interested in keeping the number of false negatives and false positives low, we opted for the  $F_1$  metric to measure the performance of the algorithms.

```
createPartition <- function(data_, p=0.7){
  # Inputs: data_ (data.frame) to split
  #         p (numeric): dataframe's proportion for train sample
  t <- unique(data_$device)
  n <- length(t)
  n.p <- round(n*p, 0)
  t.sample <- sample(t, n.p)
  train.index <- which( data_$device %in% t.sample)
  return(train.index)
}

f1 <- function (data, lev = NULL, model = NULL) {
  # Function requires to calculate F1 score within caret::train , see doc.
  precision <- posPredValue(data$pred, data$obs, positive = "Failure")
  recall <- sensitivity(data$pred, data$obs, positive = "Failure")
  f1_val <- (2 * precision * recall) / (precision + recall)
  names(f1_val) <- c("F1")
  return(f1_val)
}
```

```

set.seed(0)
data.sample$failure <- factor(data.sample$failure)
levels(data.sample$failure) <- c('NoFailure', 'Failure')
train.index <- createPartition(data.sample)
data.sample$date <- data.sample$device <- NULL
train <- data.sample[train.index, ]
test <- data.sample[-train.index, ]
fit.control <- trainControl( method = 'repeatedcv', number = 10, repeats = 3,
                             allowParallel = TRUE, classProbs = TRUE,
                             summaryFunction = f1, sampling = "down")

set.seed(0)
gbmFit1 <- train(failure ~ ., data = train, method = "gbm", trControl = fit.control,
                 verbose = FALSE)
xgb.Fit1 <- train(failure ~ ., data = train, method = "xgbTree", #tuneLength = 5, search= 'random',
                 trControl = fit.control,
                 verbose = FALSE)
rf.Fit1 <- train(failure ~ ., data = train, method = "rf", trControl = fit.control,
                verbose = FALSE)
rlg.Fit1 <- train(failure ~ ., data = train, method = "regLogistic",
                 trControl = fit.control, verbose = FALSE)

```

The test based on the Bonferroni intervals strongly suggests that XGB and RF perform better than the other methods, however, when evaluating it on the data test , we opted to only report RF's tuning results:

```

resamps <- resamples(list(GBM = gbmFit1, XGB = xgb.Fit1,
                          RF = rf.Fit1, RLG=rlg.Fit1 ))
summary(resamps)

```

```

##
## Call:
## summary.resamples(object = resamps)
##
## Models: GBM, XGB, RF, RLG
## Number of resamples: 30
##
## F1
##      Min.    1st Qu.    Median    Mean    3rd Qu.    Max. NA's
## GBM 0.01652893 0.03516260 0.04480287 0.04568529 0.05143872 0.07526882    0
## XGB 0.02419355 0.03330070 0.04371180 0.04614946 0.05542813 0.10687023    0
## RF  0.01183432 0.04102891 0.05163225 0.05286921 0.06431066 0.09600000    0
## RLG 0.02564103 0.04301242 0.05266805 0.05399095 0.05780347 0.09756098    0

```

```

summary(diff(resamps))

##
## Call:
## summary.diff.resamples(object = diff(resamps))
##
## p-value adjustment: bonferroni
## Upper diagonal: estimates of the difference
## Lower diagonal: p-value for H0: difference = 0
##
## F1
##      GBM      XGB      RF      RLG

```

```
## GBM      -0.0004642 -0.0071839 -0.0083057
## XGB 1.0000      -0.0067198 -0.0078415
## RF  0.6181 0.6889      -0.0011217
## RLG 0.1565 0.8533      1.0000
```

```
t2 <- Sys.time()
t2 - t1
```

```
## Time difference of 2.016175 mins
```

```
confusionMatrix(predict(rf.Fit1$finalModel,test), test$failure)
```

```
## Confusion Matrix and Statistics
```

```
##
##              Reference
## Prediction NoFailure Failure
## NoFailure      2349      10
## Failure         876      22
##
##              Accuracy : 0.728
##              95% CI : (0.7123, 0.7432)
##      No Information Rate : 0.9902
##      P-Value [Acc > NIR] : 1
##
##              Kappa : 0.0289
##
## Mcnemar's Test P-Value : <2e-16
##
##      Sensitivity : 0.7284
##      Specificity : 0.6875
##      Pos Pred Value : 0.9958
##      Neg Pred Value : 0.0245
##      Prevalence : 0.9902
##      Detection Rate : 0.7212
##      Detection Prevalence : 0.7243
##      Balanced Accuracy : 0.7079
##
##      'Positive' Class : NoFailure
##
```

```
confusionMatrix(predict(xgb.Fit1,test), test$failure)
```

```
## Confusion Matrix and Statistics
```

```
##
##              Reference
## Prediction NoFailure Failure
## NoFailure      2623      11
## Failure         602      21
##
##              Accuracy : 0.8118
##              95% CI : (0.7979, 0.8251)
##      No Information Rate : 0.9902
##      P-Value [Acc > NIR] : 1
##
##              Kappa : 0.0463
##
```



```

## McNemar's Test P-Value : <2e-16
##
##           Sensitivity : 0.81333
##           Specificity : 0.65625
##           Pos Pred Value : 0.99582
##           Neg Pred Value : 0.03371
##           Prevalence : 0.99018
##           Detection Rate : 0.80534
##           Detection Prevalence : 0.80872
##           Balanced Accuracy : 0.73479
##
##           'Positive' Class : NoFailure
##

t3 <- Sys.time()
set.seed(0)
tune_grid <- expand.grid(nrounds=c(100,300), max_depth = c(4:7), eta = c(0.05, 1), gamma = c(0.01),
                        colsample_bytree = c(0.75), subsample = c(0.50), min_child_weight = c(0))

xgb_fit <- train(failure ~., data = train, method = "xgbTree",
                trControl= fit.control,
                tuneGrid = tune_grid,
                tuneLength = 10)
tune_grid <- expand.grid(.mtry = (1:16))
rf_fit <- train(failure ~., data = train, method = "rf",
                trControl= fit.control,
                tuneGrid = tune_grid,
                tuneLength = 10)

t4 <- Sys.time()
t4 - t1

## Time difference of 3.655641 mins

confusionMatrix(predict(rf_fit$finalModel, test), test$failure)

## Confusion Matrix and Statistics
##
##           Reference
## Prediction NoFailure Failure
## NoFailure      2534      14
## Failure         691      18
##
##           Accuracy : 0.7835
##           95% CI : (0.769, 0.7976)
##           No Information Rate : 0.9902
##           P-Value [Acc > NIR] : 1
##
##           Kappa : 0.0304
##
## McNemar's Test P-Value : <2e-16
##
##           Sensitivity : 0.78574
##           Specificity : 0.56250
##           Pos Pred Value : 0.99451
##           Neg Pred Value : 0.02539

```

```
##           Prevalence : 0.99018
##       Detection Rate : 0.77802
##   Detection Prevalence : 0.78232
##       Balanced Accuracy : 0.67412
##
##       'Positive' Class : NoFailure
##
confusionMatrix(predict(xgb_fit, test), test$failure)
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction NoFailure Failure
## NoFailure      2048         5
## Failure        1177        27
##
##           Accuracy : 0.6371
##           95% CI : (0.6203, 0.6536)
##   No Information Rate : 0.9902
##   P-Value [Acc > NIR] : 1
##
##           Kappa : 0.025
##
## Mcnemar's Test P-Value : <2e-16
##
##           Sensitivity : 0.63504
##           Specificity : 0.84375
##           Pos Pred Value : 0.99756
##           Neg Pred Value : 0.02243
##           Prevalence : 0.99018
##           Detection Rate : 0.62880
##   Detection Prevalence : 0.63033
##       Balanced Accuracy : 0.73939
##
##       'Positive' Class : NoFailure
##
```